

Project Cardiovascular Disease Dataset

أمراض القلب والأوعية الدموية



Design & Data:

the purpose is prediction if a patient has cardiovascular disease or not, beside of patients medical information such as height, weight, age, blood pressure, glucose levels, and cholesterol levels.

the dataset is <https://www.kaggle.com/sulianova/cardiovascular-disease-dataset>

Features:

- Age | Objective Feature | age | int (days)
- Height | Objective Feature | height | int (cm) |
- Weight | Objective Feature | weight | float (kg) |
- Gender | Objective Feature | gender | categorical code |
- Systolic blood pressure | Examination Feature | ap_hi | int |
- Diastolic blood pressure | Examination Feature | ap_lo | int |
- Cholesterol | Examination Feature | cholesterol | 1: normal, 2: above normal, 3: well above normal |
- Glucose | Examination Feature | gluc | 1: normal, 2: above normal, 3: well above normal |
- Smoking | Subjective Feature | smoke | binary |
- Alcohol intake | Subjective Feature | alco | binary |
- Physical activity | Subjective Feature | active | binary |
- Presence or absence of cardiovascular disease | Target Variable | cardio | binary |

import the tools that we will need

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import precision_score, recall_score, accuracy_score, roc_auc_score
from sklearn.metrics import classification_report, confusion_matrix

%matplotlib inline

```

read the dataset file which was separated semicolon

```

In [2]: df = pd.read_csv("cardio_train.csv", sep = ';') #the data semicolon (not comma) separated
df.head(5)

```

```

Out[2]:
   id  age  gender  height  weight  ap_hi  ap_lo  cholesterol  gluc  smoke  alco  active  cardio
0  0  18393      2    168    62.0   110    80             1     1     0     0      1      0
1  1  20228      1    156    85.0   140    90             3     1     0     0      1      1
2  2  18857      1    165    64.0   130    70             3     1     0     0      0      1
3  3  17623      2    169    82.0   150   100             1     1     0     0      1      1
4  4  17474      1    156    56.0   100    60             1     1     0     0      0      0

```

check the columns names

```

In [3]: df.columns

```

```

Out[3]: Index(['id', 'age', 'gender', 'height', 'weight', 'ap_hi', 'ap_lo',
              'cholesterol', 'gluc', 'smoke', 'alco', 'active', 'cardio'],
              dtype='object')

```

and check the the number of records and features

```

In [4]: df.shape

```

```

Out[4]: (70000, 13)

```

```

In [5]: #check for null
df.isnull().sum().any()

```

```

Out[5]: False

```

see descriptive statistics of the features

```

In [6]: df.describe()

```

```

Out[6]:
   id  age  gender  height  weight  ap_hi  ap_lo
count 70000.000000 70000.000000 70000.000000 70000.000000 70000.000000 70000.000000 70000.000000
mean  49972.419900 19468.865814      1.349571  164.359229    74.205690   128.817286    96.63041
std   28851.302323  2467.251667    0.476838    8.210126    14.395757   154.011419   188.47253

```

| | id | age | gender | height | weight | ap_hi | ap_lo |
|------------|--------------|--------------|----------|------------|------------|--------------|--------------|
| min | 0.000000 | 10798.000000 | 1.000000 | 55.000000 | 10.000000 | -150.000000 | -70.000000 |
| 25% | 25006.750000 | 17664.000000 | 1.000000 | 159.000000 | 65.000000 | 120.000000 | 80.000000 |
| 50% | 50001.500000 | 19703.000000 | 1.000000 | 165.000000 | 72.000000 | 120.000000 | 80.000000 |
| 75% | 74889.250000 | 21327.000000 | 2.000000 | 170.000000 | 82.000000 | 140.000000 | 90.000000 |
| max | 99999.000000 | 23713.000000 | 2.000000 | 250.000000 | 200.000000 | 16020.000000 | 11000.000000 |

EDA:

cleaning:

- 1 drop column `id` as it is irrelevant to target variable.
- 2 Transform `age` column into years instead of days.
- 3 clean the unrealistic values in `ap_hi` and `ap_lo`
- 4 clean the unrealistic values in `height` and `weight`
- 5 add **Body Mass Index (BMI)** column
- 6 change the encode for `gender` from 1-2 to F-M
- 7 check for **duplicated** data

1 drop column `id` as it is irrelevant to target variable.

```
In [7]: #1 columns before drop id
df.shape
```

```
Out[7]: (70000, 13)
```

```
In [8]: #1 drop id
df=df.drop('id',axis=1)
```

```
In [9]: #1 columns after drop id
df.shape
```

```
Out[9]: (70000, 12)
```

2 Transform `age` column into years instead of days.

```
In [10]: #2 before
df.age[:3]
```

```
Out[10]: 0    18393
         1    20228
```

```
2    18857
Name: age, dtype: int64
```

```
In [11]: #2 age from days to years
df.age = np.round(df.age/365.25, decimals=1)
```

use round function from numpy for rounding decimals


```
In [12]: #2 after
df[['age']].sort_values(by=['age'], ascending=True)
```

```
Out[12]:
```

| | age |
|-------|------|
| 22343 | 29.6 |
| 30666 | 29.7 |
| 6219 | 29.8 |
| 55905 | 30.0 |
| 44857 | 39.1 |
| ... | ... |
| 57191 | 64.9 |
| 50714 | 64.9 |
| 20931 | 64.9 |
| 36603 | 64.9 |
| 68005 | 64.9 |

70000 rows × 1 columns

3 clean the unrealistic values in ap_hi and ap_lo



the normal human blood pressure in Systolic blood pressure (ap_hi) between 70-190 and in Diastolic blood pressure (ap_lo) between 40-120

we have unreal readings in our data and we will drop them

```
In [13]: #3 before  
df[['ap_hi']].sort_values(by=['ap_hi'],ascending=True)
```

```
Out[13]:
```

| | ap_hi |
|-------|-------|
| 35040 | -150 |
| 23988 | -140 |
| 46627 | -120 |
| 25240 | -120 |
| 16021 | -115 |
| ... | ... |
| 47253 | 14020 |
| 25464 | 14020 |
| 25519 | 14020 |
| 46912 | 14020 |
| 40852 | 16020 |

70000 rows × 1 columns

```
In [14]: #3 before  
df[['ap_lo']].sort_values(by=['ap_lo'],ascending=True)
```

```
Out[14]:
```

| | ap_lo |
|-------|-------|
| 60106 | -70 |
| 40330 | 0 |
| 42397 | 0 |
| 56950 | 0 |
| 63787 | 0 |
| ... | ... |
| 43434 | 9800 |
| 68538 | 10000 |
| 23849 | 10000 |
| 2381 | 10000 |
| 43326 | 11000 |

70000 rows × 1 columns

test if the record want to drop is less than 5% of the data

```
In [15]: df.shape[0:1] #number of rows
```

```
Out[15]: (70000,)
```

5% of 70000 is 3500

so, more than 66500 is ok

```
In [16]: # check the rows numbers if we drop  
 #(ap_hi) between 70-190  
 #(ap_lo) between 40-120  
del_df= df[(df['ap_lo']<=120) & (df['ap_hi']<=190)]  
del_df= df[(df['ap_lo']>=40) & (df['ap_hi']>=70)]  
del_df.shape[0:1]
```

```
Out[16]: (69757,)
```

its ok

```
In [17]: #3 drop unreal readings  
 #(ap_hi) between 70-190  
 #(ap_lo) between 40-120  
df= df[(df['ap_lo']<=120) & (df['ap_hi']<=190)]  
df= df[(df['ap_lo']>=40) & (df['ap_hi']>=70)]
```

now the max and min realistic

```
In [18]: df[['ap_lo', 'ap_hi']].describe()
```

```
Out[18]:
```

| | ap_lo | ap_hi |
|--------------|--------------|--------------|
| count | 68548.000000 | 68548.000000 |
| mean | 81.247593 | 126.452179 |
| std | 9.313123 | 16.301674 |
| min | 40.000000 | 70.000000 |
| 25% | 80.000000 | 120.000000 |
| 50% | 80.000000 | 120.000000 |
| 75% | 90.000000 | 140.000000 |
| max | 120.000000 | 190.000000 |

4 clean the unrealistic values in height and weight

| Adults Weight to Height Ratio Chart | | |
|-------------------------------------|----------------------------------|----------------------------------|
| Height - Ft. In. (cms) | Female | Male |
| 4' 6" - (137 cm) | 63 - 77 lb - (28.5 - 34.9 kg) | 63 - 77 lb - (28.5 - 34.9 kg) |
| 4' 7" - (140 cm) | 68 - 83 lb - (30.8 - 37.6 kg) | 68 - 84 lb - (30.8 - 38.1 kg) |
| 4' 8" - (142 cm) | 72 - 88 lb - (32.6 - 39.9 kg) | 74 - 90 lb - (33.5 - 40.8 kg) |
| 4' 9" - (145 cm) | 77 - 94 lb - (34.9 - 42.6 kg) | 79 - 97 lb - (35.8 - 43.9 kg) |
| 4' 10" - (147 cm) | 81 - 99 lb - (36.4 - 44.9 kg) | 85 - 103 lb - (38.5 - 46.7 kg) |
| 4' 11" - (150 cm) | 86 - 105 lb - (39 - 47.6 kg) | 90 - 110 lb - (40.8 - 49.9 kg) |
| 5' 0" - (152 cm) | 90 - 110 lb - (40.8 - 49.9 kg) | 95 - 117 lb - (43.1 - 53 kg) |
| 5' 1" - (155 cm) | 95 - 116 lb - (43.1 - 52.6 kg) | 101 - 123 lb - (45.8 - 55.8 kg) |
| 5' 2" - (157 cm) | 99 - 121 lb - (44.9 - 54.9 kg) | 106 - 130 lb - (48.1 - 58.9 kg) |
| 5' 3" - (160 cm) | 104 - 127 lb - (47.2 - 57.6 kg) | 112 - 136 lb - (50.8 - 61.6 kg) |
| 5' 4" - (163 cm) | 108 - 132 lb - (49 - 59.9 kg) | 117 - 143 lb - (53 - 64.8 kg) |
| 5' 5" - (165 cm) | 113 - 138 lb - (51.2 - 62.6 kg) | 122 - 150 lb - (55.3 - 68 kg) |
| 5' 6" - (168 cm) | 117 - 143 lb - (53 - 64.8 kg) | 128 - 156 lb - (58 - 70.7 kg) |
| 5' 7" - (170 cm) | 122 - 149 lb - (55.3 - 67.6 kg) | 133 - 163 lb - (60.3 - 73.9 kg) |
| 5' 8" - (173 cm) | 126 - 154 lb - (57.1 - 69.8 kg) | 139 - 169 lb - (63 - 76.6 kg) |
| 5' 9" - (175 cm) | 131 - 160 lb - (59.4 - 72.6 kg) | 144 - 176 lb - (65.3 - 79.8 kg) |
| 5' 10" - (178 cm) | 135 - 165 lb - (61.2 - 74.8 kg) | 149 - 183 lb - (67.6 - 83 kg) |
| 5' 11" - (180 cm) | 140 - 171 lb - (63.5 - 77.5 kg) | 155 - 189 lb - (70.3 - 85.7 kg) |
| 6' 0" - (183 cm) | 144 - 176 lb - (65.3 - 79.8 kg) | 160 - 196 lb - (72.6 - 88.9 kg) |
| 6' 1" - (185 cm) | 149 - 182 lb - (67.6 - 82.5 kg) | 166 - 202 lb - (75.3 - 91.6 kg) |
| 6' 2" - (188 cm) | 153 - 187 lb - (69.4 - 84.8 kg) | 171 - 209 lb - (77.5 - 94.8 kg) |
| 6' 3" - (191 cm) | 158 - 193 lb - (71.6 - 87.5 kg) | 176 - 216 lb - (79.8 - 98 kg) |
| 6' 4" - (193 cm) | 162 - 198 lb - (73.5 - 89.8 kg) | 182 - 222 lb - (82.5 - 100.6 kg) |
| 6' 5" - (195 cm) | 167 - 204 lb - (75.7 - 92.5 kg) | 187 - 229 lb - (84.8 - 103.8 kg) |
| 6' 6" - (198 cm) | 171 - 209 lb - (77.5 - 94.8 kg) | 193 - 235 lb - (87.5 - 106.5 kg) |
| 6' 7" - (201 cm) | 176 - 215 lb - (79.8 - 97.5 kg) | 198 - 242 lb - (89.8 - 109.7 kg) |
| 6' 8" - (203 cm) | 180 - 220 lb - (81.6 - 99.8 kg) | 203 - 249 lb - (92 - 112.9 kg) |
| 6' 9" - (205 cm) | 185 - 226 lb - (83.9 - 102.5 kg) | 209 - 255 lb - (94.8 - 115.6 kg) |
| 6' 10" - (208 cm) | 189 - 231 lb - (85.7 - 104.8 kg) | 214 - 262 lb - (97 - 118.8 kg) |
| 6' 11" - (210 cm) | 194 - 237 lb - (88 - 107.5 kg) | 220 - 268 lb - (99.8 - 121.5 kg) |
| 7' 0" - (213 cm) | 198 - 242 lb - (89.8 - 109.7 kg) | 225 - 275 lb - (102 - 124.7 kg) |

adlut height between 137-213

adlut weight between 28.5-124.7

```
In [19]: # the max and min not realiztic
df[['height','weight']].describe()
```

```
Out[19]:
```

| | height | weight |
|--------------|--------------|--------------|
| count | 68548.000000 | 68548.000000 |
| mean | 164.362111 | 74.089397 |
| std | 8.179812 | 14.304198 |
| min | 55.000000 | 11.000000 |
| 25% | 159.000000 | 65.000000 |
| 50% | 165.000000 | 72.000000 |
| 75% | 170.000000 | 82.000000 |
| max | 250.000000 | 200.000000 |

```
In [20]: # check the rows numbers if we drop
del_df = df[(df['height']<=213.0) & (df['height']>=137.0)]
del_df = df[(df['weight']<=124.7) & (df['weight']>=28.5)]
del_df.shape[0:1]
```

Out[20]: (68185,)

its less than 5% so its ok

```
In [21]: #4 drop the unreals
#adlut height between 137-213
#adlut weight between 28.5-124.7
df = df[(df['height']<=213.0) & (df['height']>=137.0)]
df = df[(df['weight']<=124.7) & (df['weight']>=28.5)]
```

```
In [22]: # the max and min realiztic
df[['height','weight']].describe()
```

```
Out[22]:
```

| | height | weight |
|--------------|--------------|--------------|
| count | 68072.000000 | 68072.000000 |
| mean | 164.427562 | 73.766033 |
| std | 7.799930 | 13.536602 |
| min | 137.000000 | 29.000000 |
| 25% | 159.000000 | 65.000000 |
| 50% | 165.000000 | 72.000000 |
| 75% | 170.000000 | 82.000000 |
| max | 207.000000 | 124.000000 |

5 add Body Mass Index (BMI) column

Let's create a new feature - Body Mass Index (BMI):

$$BMI = \frac{mass_{kg}}{height_m^2},$$

it's easier to compute BMI instead of height and weight

| BMI | Weight status |
|------------|-------------------|
| Below 18.5 | Underweight |
| 18.5-24.9 | Normal weight |
| 25.0-29.9 | Overweight |
| 30.0-34.9 | Obesity class I |
| 35.0-39.9 | Obesity class II |
| Above 40 | Obesity class III |

BMI btween 18.5-40

```
In [23]: # add the feature
df['BMI'] = df['weight']/((df['height']/100)**2)
```

```
In [24]: # not realiztic
df[['BMI']].sort_values(by=['BMI'],ascending=True)
```

Out[24]:

| | BMI |
|--|-----|
|--|-----|

| | |
|--------------|-----------|
| 60699 | 9.917581 |
| 16906 | 10.726644 |
| 18559 | 11.718750 |
| 58200 | 12.254473 |
| 16322 | 12.855831 |
| ... | ... |
| 66997 | 54.666667 |
| 69708 | 55.459105 |
| 15319 | 56.295740 |
| 49377 | 57.870370 |
| 28448 | 58.024202 |

68072 rows × 1 columns

```
In [25]: # check the raws numbers if we drop
del_df = df[(df['BMI']<40.0) & (df['BMI']>18.5)]
del_df.shape[0:1]
```

Out[25]: (66054,)

its less than 5% so its ok

```
In [26]: # cant be less than 18.5 and more than 40, so drop unrels
df = df[(df['BMI']<40.0) & (df['BMI']>18.5)]
```

```
In [27]: # realiztic
df[['BMI']].sort_values(by=['BMI'],ascending=True)
```

Out[27]:

| | BMI |
|--|-----|
|--|-----|

| | |
|--------------|-----------|
| 46646 | 18.507766 |
| 25855 | 18.507766 |
| 16208 | 18.507766 |

| | BMI |
|-------|-----------|
| 60844 | 18.507766 |
| 10612 | 18.507766 |
| ... | ... |
| 8959 | 39.958377 |
| 65510 | 39.958377 |
| 35444 | 39.959508 |
| 56331 | 39.965649 |
| 52820 | 39.965649 |

66054 rows × 1 columns

6 change the encode for gender from 1-2 to F-M

```
In [28]: df["gender"].unique()
```

```
Out[28]: array([2, 1], dtype=int64)
```

we dont have info who is 1 or 2

so this next cell can know besides on info that the height avg. of male is more than female

```
In [29]: #take the mean of both 1 and 2
a = df[df["gender"]==1]["height"].mean()
b = df[df["gender"]==2]["height"].mean()

# compare
if a > b:
    gender = "male"
    gender1 = "female"
else:
    gender = "female"
    gender1 = "male"

print("Gender:1 is " + gender + " & Gender:2 is " + gender1)
```

Gender:1 is female & Gender:2 is male

```
In [30]: # use lambda fun. to change the 1,2 to f,m
df['gender'] = df['gender'].apply(lambda x: 'F' if x == 1 else 'M')
```

```
In [31]: df.gender.value_counts()
```

```
Out[31]: F    42665
M    23389
Name: gender, dtype: int64
```

7 check for duplicated data

```
In [32]: # check the rows and columns  
df.shape
```

```
Out[32]: (66054, 13)
```

```
In [33]: # see how much duplicated rows  
df.duplicated().sum()
```

```
Out[33]: 673
```

```
In [34]: #drpp the 673 duplicated raw  
df.drop_duplicates(inplace=True)
```

```
In [35]: # now we should have 0  
df.duplicated().sum()
```

```
Out[35]: 0
```

```
In [36]: # check the rows and columns  
df.shape
```

```
Out[36]: (65381, 13)
```


EDA:

Questions:

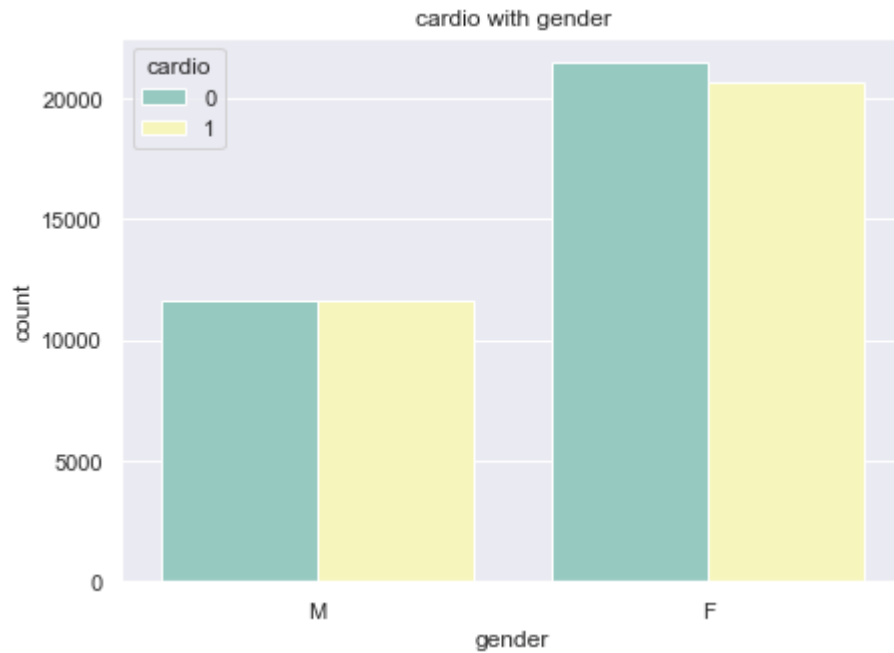
- 1 Which gender has more has cardiovascular disease?
- 2 Which age group has more has cardiovascular disease?
- 3 Blood pressure effect the cardiovascular disease?

1 Which gender has more cardiovascular disease?

```
In [67]: sns.set_style('whitegrid') #to make white grid in the graph  
sns.set(rc={'figure.figsize':(7,5)}) # for edit size  
sns.countplot(df.gender,hue=df.cardio, palette="Set3"); #make the graph  
plt.title('cardio with gender'); # add title
```

C:\Users\ichra\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will

result in an error or misinterpretation.
warnings.warn(



```
In [39]: # show count of cardio in every gender
df.groupby(['gender', 'cardio']).agg({'cardio': 'count'})
```

```
Out[39]:
```

| | | cardio | |
|--------|--------|--------|--|
| gender | cardio | | |
| F | 0 | 21467 | |
| | 1 | 20640 | |
| M | 0 | 11657 | |
| | 1 | 11617 | |

the females is more than males so lets try another way with percent (%)

```
In [40]: # show count of cardio in every gender in percent
gender_cardio = df.groupby(['gender', 'cardio']).agg({'cardio': 'count'})
gender = df.groupby(['gender']).agg({'cardio': 'count'})
gender_cardio.div(gender, level='gender') * 100
```

```
Out[40]:
```

| | | cardio | |
|--------|--------|-----------|--|
| gender | cardio | | |
| F | 0 | 50.982022 | |
| | 1 | 49.017978 | |
| M | 0 | 50.085933 | |
| | 1 | 49.914067 | |


the females more than males in this dataset but when it come with percent they almost **50/50** to

have cardiovascular disease

2 Which age group has more has cardiovascular disease?

most the ages that has cardio between 50-60

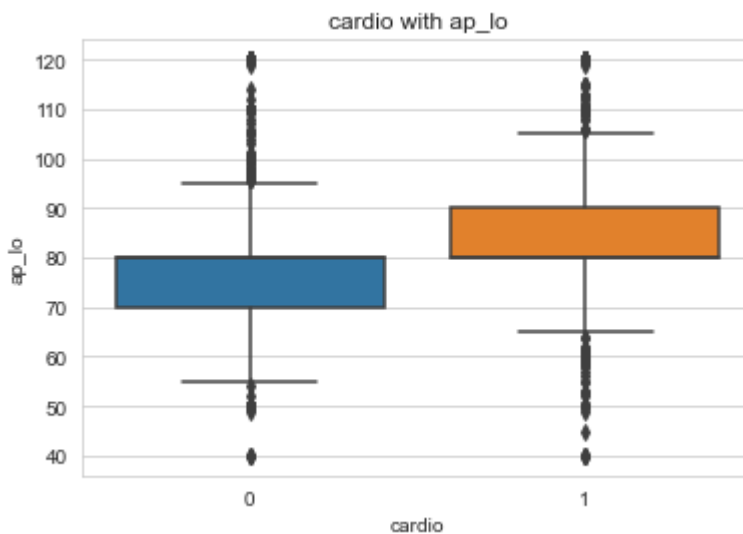
```
In [41]: sns.boxplot( x="cardio", y="age", data=df); # make boxplot  
plt.title('cardio with age'); #set title
```



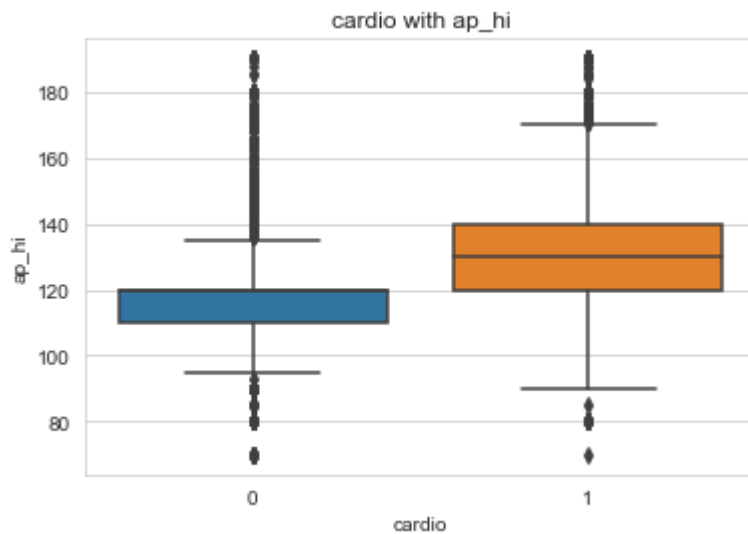
3 Patients with cardiovascular disease have high Blood pressure?

Yes, they have higher than patients they dont

```
In [42]: sns.boxplot( x="cardio", y="ap_lo", data=df); # make boxplot  
plt.title('cardio with ap_lo'); #set title
```



```
In [43]: sns.boxplot( x="cardio", y="ap_hi", data=df); # make boxplot  
plt.title('cardio with ap_hi'); #set title
```



Model :

we will use **Logistic Regression** for predictions if a patient has cardiovascular disease or not

- Step 1: Feature selection
- Step 2: Split the data
- Step 3: Scale the train data
- Step 4: Use Logistic Regression

Step 1: Feature selection

```
In [45]: df.corr()['cardio'].sort_values(ascending=False) #show correlation for cardio with othe
```

```
Out[45]: cardio          1.000000
ap_hi          0.429374
ap_lo          0.336169
age            0.239377
cholesterol    0.218216
BMI            0.183562
weight         0.167538
gluc           0.086210
height        -0.007635
alco           -0.010798
smoke          -0.017629
active         -0.034054
Name: cardio, dtype: float64
```

Features ap_hi, ap_lo are higher (moderate positive correlation) correlation with cardio

Feature cardio is the predict target to predict the patient have cardiovascular disease or not

Step 2 : Split the data

```
In [46]: X = df[['ap_hi', 'ap_lo']] #seleced feature
        y = df['cardio'] #target feature
```

```
In [47]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=4
```

Step 3: Scale the train data and test data

```
In [48]: sc = StandardScaler() # use StandardScaler
        lr = LogisticRegression() # use LogisticRegression
```

```
In [49]: X_train_sc = sc.fit_transform(X_train) # Scale train
```

```
In [50]: X_test_sc = sc.transform(X_test) # Scale test
```

Step 4: Use Logistic Regression

```
In [51]: lr.fit(X_train_sc,y_train) # use LogisticRegression
```

```
Out[51]: LogisticRegression()
```

```
In [52]: lr.score(X_train_sc,y_train) #show the train score
```

```
Out[52]: 0.7125458855919241
```

train score = 0.71

```
In [53]: y_pred=lr.predict(X_test_sc) # Estimated target
```

```
In [54]: accuracy_score(y_test,y_pred) #test result
```

```
Out[54]: 0.7056664372562514
```

the train score is 0.71

and the test score was 0.705 almost 0.71

```
In [55]: df['cardio'].value_counts() # balanced data
```

```
Out[55]: 0    33124
        1    32257
        Name: cardio, dtype: int64
```

we have balanced data

Classification report

to measure the quality of predictions

```
In [57]: print(classification_report(y_test,y_pred))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.67 | 0.81 | 0.73 | 6517 |
| 1 | 0.76 | 0.61 | 0.67 | 6560 |
| accuracy | | | 0.71 | 13077 |
| macro avg | 0.71 | 0.71 | 0.70 | 13077 |
| weighted avg | 0.71 | 0.71 | 0.70 | 13077 |

Precision is Accuracy of positive predictions.

```
In [58]: print(confusion_matrix(y_test, y_pred))
```

```
[[5255 1262]
 [2587 3973]]
```

| | | Predicted class | |
|--------------|----------|----------------------|----------------------|
| | | <i>P</i> | <i>N</i> |
| Actual Class | <i>P</i> | True Positives (TP) | False Negatives (FN) |
| | <i>N</i> | False Positives (FP) | True Negatives (TN) |

```
In [ ]:
```